


# inclass07

## Clustering

First, lets make up some data to cluster so we can get a feel for these methods and how to work with them.

We can use the “rnorm” function to get random numbers from a normal distribution around a given “mean”

```
hist(rnorm(50000, mean=3))
```



Lets gets 30 points with a mean of 3

```
temp <- c(rnorm(30, mean=3),  
rnorm(30, mean=-3))
```


put two of these together

```
x <- cbind(x=temp, y=rev(temp))  
x
```

	x	y
[1,]	2.742507	-2.237313
[2,]	1.930372	-3.421799
[3,]	3.697294	-1.977868
[4,]	4.056811	-1.331041
[5,]	2.558193	-2.660298
[6,]	1.454726	-3.662458
[7,]	4.301499	-3.320913
[8,]	4.405483	-3.426306
[9,]	2.270497	-1.374508
[10,]	2.380468	-3.277130
[11,]	4.520089	-2.603031
[12,]	1.789649	-0.324164
[13,]	3.303600	-1.313204
[14,]	2.726551	-4.382811
[15,]	4.030920	-2.133516
[16,]	1.724160	-2.337513
[17,]	1.450845	-5.106174
[18,]	2.517830	-4.450796
[19,]	2.274701	-4.279307
[20,]	2.017993	-2.820837
[21,]	3.858287	-3.503523
[22,]	2.914180	-3.539410
[23,]	2.975645	-3.815064
[24,]	3.456224	-2.686058
[25,]	2.861493	-1.952562
[26,]	1.959717	-2.000008
[27,]	3.734885	-1.754357
[28,]	2.525061	-1.055127
[29,]	1.903224	-2.441048
[30,]	1.838105	-3.737298
[31,]	-3.737298	1.838105
[32,]	-2.441048	1.903224
[33,]	-1.055127	2.525061

[34,]	-1.754357	3.734885
[35,]	-2.000008	1.959717
[36,]	-1.952562	2.861493
[37,]	-2.686058	3.456224
[38,]	-3.815064	2.975645
[39,]	-3.539410	2.914180
[40,]	-3.503523	3.858287
[41,]	-2.820837	2.017993
[42,]	-4.279307	2.274701
[43,]	-4.450796	2.517830
[44,]	-5.106174	1.450845
[45,]	-2.337513	1.724160
[46,]	-2.133516	4.030920
[47,]	-4.382811	2.726551
[48,]	-1.313204	3.303600
[49,]	-0.324164	1.789649
[50,]	-2.603031	4.520089
[51,]	-3.277130	2.380468
[52,]	-1.374508	2.270497
[53,]	-3.426306	4.405483
[54,]	-3.320913	4.301499
[55,]	-3.662458	1.454726
[56,]	-2.660298	2.558193
[57,]	-1.331041	4.056811
[58,]	-1.977868	3.697294
[59,]	-3.421799	1.930372
[60,]	-2.237313	2.742507

```
plot(x)
```



##K-means clustering

Very popular clustering methods, especially for big datasets, that we can use with the “kmeans()” function in base R.

```
km <- kmeans(x, centers=2)
km
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

	x	y
1	-2.764181	2.806034
2	2.806034	-2.764181

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 62.53017 62.53017
(between_SS / total_SS = 88.2 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

```
kmeans(x, centers=2, nstart=20)
```

K-means clustering with 2 clusters of sizes 30, 30

Cluster means:

```
      x      y
1 -2.764181  2.806034
2  2.806034 -2.764181
```

Clustering vector:

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

Within cluster sum of squares by cluster:

```
[1] 62.53017 62.53017
(between_SS / total_SS =  88.2 %)
```

Available components:

```
[1] "cluster"      "centers"      "totss"        "withinss"     "tot.withinss"
[6] "betweenss"    "size"         "iter"         "ifault"
```

Q: What component of your result objects details

-cluster size?

```
km$size
```

```
[1] 30 30
```

-cluster assignment/membership?

```
km$cluster
```

```
[1] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 1 1 1 1 1 1 1 1
[39] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
```

cluster center?


```
km$centers
```

```
      x      y
1 -2.764181 2.806034
2  2.806034 -2.764181
```

Q. Plot x colored by the kmeans cluster assignment and add cluster points as blue points


#pch made it a square, cex made the square bigger

```
plot(x, col=km$cluster)
points(km$centers, col="blue", pch=15, cex=3)
```



Q. Lets cluster into 3 groups or same “x” data and make a plot

```
km <- kmeans(x,3)
plot(x, col=km$cluster)
```



#Hierarchical clustering

we can use the 'hclust()' function for Hierarchical Clustering. Unlike 'kmeans()' where we could just pass in our data as input, we need to give 'hclust()' a distance matrix.

We will use the 'dist()' function to start with

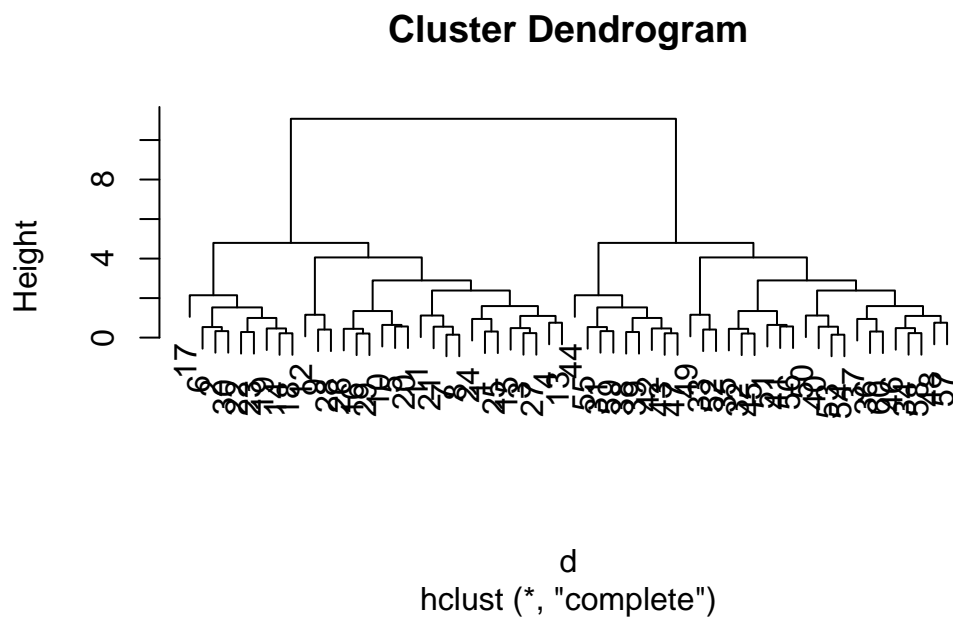
```
d <- dist(x)
hc <- hclust(d)
hc
```

Call:

```
hclust(d = d)
```

```
Cluster method : complete
Distance       : euclidean
Number of objects: 60
```

```
plot(hc)
```




I can now ‘cut’ my tree with “cutree()” to yield a cluster membership vector

```
grps <- cutree(hc,h=8)
grps
```

```
[1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2
[39] 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2 2
```

```
plot(x, col=grps)
```





You can also tell “`cutree()`” to cut where it yields k groups

```
cutree(hc, k=2)
```

[illegible]

## #Principal Component Analysis (PCA)

```
url <- "https://tinyurl.com/UK-foods"
x <- read.csv(url)
```

Q1. How many rows and columns are in your new data frame named x? What R functions could you use to answer this questions?

find the number of rows and columns

 $\dim(\mathbf{x})$ 

[1] 17 5

Preview the first 6 rows

```
head(x)
```

	X	England	Wales	Scotland	N.Ireland
1	Cheese	105	103	103	66
2	Carcass_meat	245	227	242	267
3	Other_meat	685	803	750	586
4	Fish	147	160	122	93
5	Fats_and_oils	193	235	184	209
6	Sugars	156	175	147	139

fix the labeling of the rows, they are set as the first column of our x dataframe instead of proper row names

```
rownames(x) <- x[,1]  
x <- x[,-1]  
head(x)
```

	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267
Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

```
#now re check number of rows and columns to show it's corrected  
dim(x)
```

```
[1] 17  4
```

```
x <- read.csv(url, row.names=1)  
head(x)
```


	England	Wales	Scotland	N.Ireland
Cheese	105	103	103	66
Carcass_meat	245	227	242	267

Other_meat	685	803	750	586
Fish	147	160	122	93
Fats_and_oils	193	235	184	209
Sugars	156	175	147	139

Q2. Which approach to solving the ‘row-names problem’ mentioned above do you prefer and why? Is one approach more robust than another under certain circumstances?

the second is preferred because if you run the code `x<-x[,-1]` multiple times, it will keep shifting the columns

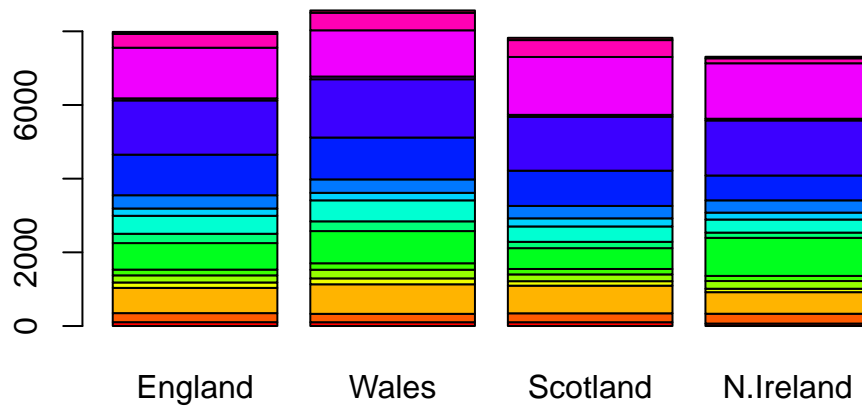
```
barplot(as.matrix(x), beside=T, col=rainbow(nrow(x)))
```



Q3: Changing what optional argument in the above `barplot()` function results in the following plot?

can change the `beside` factor to make it stacked


```
barplot(as.matrix(x), beside=F, col=rainbow(nrow(x)))
```



Q5: Generating all pairwise plots may help somewhat. Can you make sense of the following code and resulting figure? What does it mean if a given point lies on the diagonal for a given plot?

This figure shows all countries compared separately in separate plots, and you can match up their x and y axis. If a given point lies on the diagonal for a given plot, it means that at that point, both countries have the same value/correlate.

```
pairs(x, col=rainbow(10), pch=16)
```



Q6. What is the main differences between N. Ireland and the other countries of the UK in terms of this data-set?

Northern Ireland has lower values on average than other countries of the UK

#Lets use PCA to make this plot easier


```
pca <- prcomp( t(x) )
summary(pca)
```

Importance of components:

	PC1	PC2	PC3	PC4
Standard deviation	324.1502	212.7478	73.87622	4.189e-14
Proportion of Variance	0.6744	0.2905	0.03503	0.000e+00
Cumulative Proportion	0.6744	0.9650	1.00000	1.000e+00

Q7. Complete the code below to generate a plot of PC1 vs PC2. The second line adds text labels over the data points.

```
plot(pca$x[,1], pca$x[,2], xlab="PC1", ylab="PC2", xlim=c(-270,500))
text(pca$x[,1], pca$x[,2], colnames(x))
```



add colors

```
plot(pca$x[,1], pca$x[,2], col= c("orange", "red", "blue", "green"), xlab="PC1", ylab="PC2")
text(pca$x[,1], pca$x[,2], colnames(x))
```

